

TELEMETRY ONBOARDING

Roberto A. Vitillo, Mark Reid

1. What is Unified Telemetry?
2. Executive dashboards
3. Telemetry dashboards
4. Adding a probe
5. What is a ping
6. Data pipeline
7. Experiments
8. Offline processing
9. Stats
10. Privacy Policy
11. Data Collection Policy

WHAT IS UNIFIED TELEMETRY?

- Measures how Firefox behaves in the real world [1].
- Collects non-personal information about performance, hardware, usage, errors, and customizations
- The data, used in aggregate,
 - ▶ allows to identify new issues and regressions
 - ▶ allows to conduct longitudinal studies and experiments

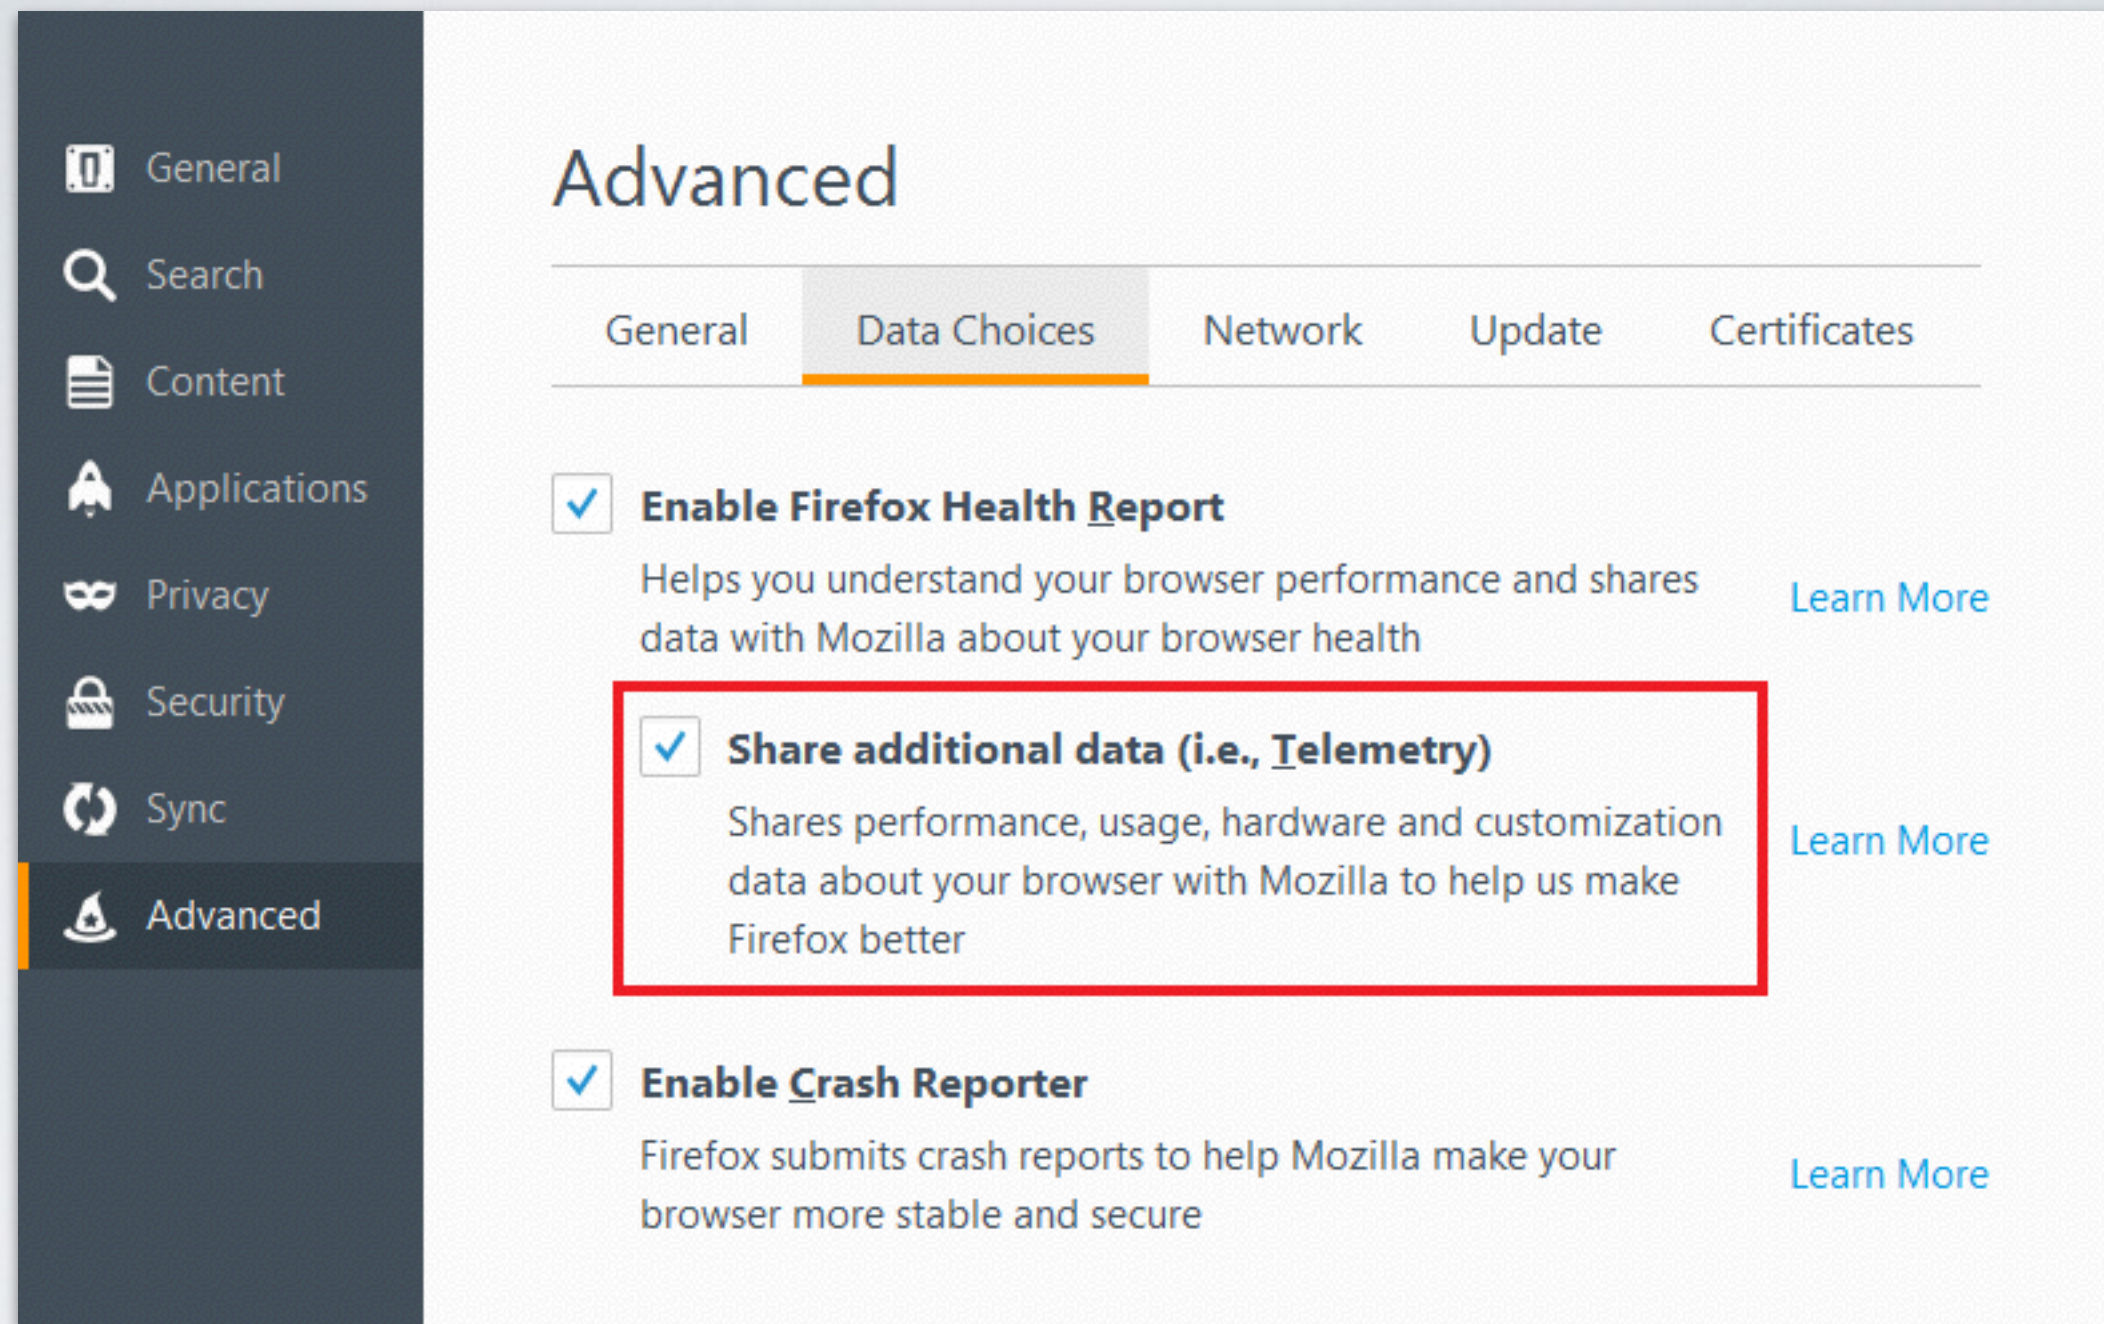
[1] https://wiki.mozilla.org/Unified_Telemetry

UNIFIED?

- Classic Telemetry [1], designed to send performance and usage data
 - ▶ enabled by default on pre-release channels and disabled by default on release
- Firefox Health Report [2] is a data source intended to be representative and support longitudinal study.
 - ▶ enabled by default on release and pre-release

[1] <https://wiki.mozilla.org/Telemetry>

[2] https://wiki.mozilla.org/Firefox_Health_Report



Add the following definitions to your mozconfig file:

```
export MOZILLA_OFFICIAL=1
export MOZ_TELEMETRY_REPORTING=1
```

Telemetry Data

about:telemetry

Search

Telemetry Data

This page shows the information about performance, hardware, usage and customisations collected by Telemetry. This information is submitted to Mozilla to help improve Mozilla Firefox.

FHR data upload is **enabled**. [Change](#)

Extended Telemetry recording is **enabled**. [Change](#)

Ping data source:

☒ Current ping data

☐ Archived ping data

☐ Show subsession data

General Data

Click to toggle section

Environment Data

Click to toggle section

Telemetry Log

Click to toggle section

Slow SQL Statements

(No data collected)

Browser Hangs

Click to toggle section

Thread Hangs

(No data collected)

Histograms

Click to toggle section

Keyed Histograms

Click to toggle section

Simple Measurements

Click to toggle section

Late Writes

Click to toggle section

System Information

Click to toggle section

Add-on Details

Click to toggle section

Histograms Collected by Add-ons

(No data collected)

REFERENCES

- https://wiki.mozilla.org/Unified_Telemetry
- <https://wiki.mozilla.org/Telemetry>
- https://wiki.mozilla.org/Firefox_Health_Report

EXECUTIVE DASHBOARDS

CONFIDENTIAL — Only for staff and contributors under NDA — Do not share

Mozilla Summary Dashboards


<https://metrics.services.mozilla.com/>

TELEMETRY DASHBOARDS

Histogram Dashboard V4

[Report bug](#)[Evolution View](#)[Dashboards Home](#)[Usage Tutorial](#)

SIMPLE MEASURES UPTIME distribution for Firefox Desktop nightly 43, on any OS (67) any architecture (3) with any e10s setting (2) as any process type (2) and compare by none



Histogram Type	exponential
Metric Count	4.92M
Sample Count	4.92M
Sample Sum	1.26B
Number of dates	43
Selected Dates	2015/08/10 to 2015/09/21

5th Percentile	0
25th Percentile	2.2
Median	13.07
75th Percentile	66.67
95th Percentile	573.68

Notice percentiles are estimated based on values in the histogram. Values are only guaranteed to be accurate to the nearest bucket.


[Export CSV](#)[Export JSON](#)[Advanced Settings](#)

Evolution Dashboard V4

[Report bug](#)[Distribution View](#)[Dashboards Home](#)[Usage Tutorial](#)

Median GC_MS from **nightly 41** to **nightly 43** for **Firefox Desktop** on **any OS (72)** **any architecture (3)** with **any e10s setting (2)** as **any process type (2)**

Time spent running JS GC (ms)



<https://telemetry.mozilla.org/new-pipeline/evo.html>

Aggregates can be accessed through a javascript library...

Examples

Getting versions in a range:

```
Telemetry.init(function() {  
    var versions = Telemetry.getVersions("nightly/40", "nightly/42");  
    console.log("Versions between nightly 40 to nightly 42 (inclusive):\n" + versions.join("\n"));  
});
```

Getting a list of measures:

```
Telemetry.init(function() {  
    Telemetry.getFilterOptions("nightly", "42", function(filterOptions) {  
        console.log("Available measures:\n" + filterOptions.metric.join("\n"));  
    });  
});
```

Getting the dates for which there are submissions for GC_MS on nightly 42 on Windows:

```
Telemetry.init(function() {  
    Telemetry.getEvolution("nightly", "42", "GC_MS", {os: "Windows_NT"}, true, function(evolutionMap) {  
        console.log("Available dates:\n" + evolutionMap[""].dates().join("\n"));  
    });  
});
```


...or directly using the HTTP endpoints

API

Aggregates are made available through a HTTP API. There are two kinds of aggregates: per submission date and per build-id. To access the aggregates use the `aggregates_by/build_id/` and `aggregates_by/submission_date/` prefix respectively.

The following examples are based on build-id aggregates.

Get available channels:

```
curl -X GET http://SERVICE/aggregates_by/build_id/channels/  
["nightly","beta","aurora"]
```

Get a list of options for the available dimensions on a given channel and version:

```
curl -X GET "http://SERVICE/filters/?channel=nightly&version=42"  
{  
  "metric":["A11Y_CONSUMERS","A11Y_IATABLE_USAGE_FLAG",...],  
  "application":["Fennec","Firefox"],  
  ...  
}
```

Get a list of available build-ids for a given channel:

```
curl -X GET "http://SERVICE/aggregates_by/build_id/channels/nightly/dates/"  
[{"date":"20150630","version":"42"}, {"date":"20150629","version":"42"}]
```

https://github.com/vitillo/python_mozagggregator/blob/master/README.md

For example: https://aggregates.telemetry.mozilla.org/aggregates_by/submission_date/channels/nightly/?version=42&dates=20150709&metric=GC_MS&application=Firefox&child=false

Telemetry alerts

Login

Histogram Regression Detector

Telemetry Main-thread IO

Unpacked Add-on File Scan

Filter Metrics By:

Select Metric:

From:


2015-08-22

To:

2015-09-22

Failure to load alerts, reason: Not Found

MEMORY_HEAP_COMMITTED_UNUSED_RATIO



<http://alerts.telemetry.mozilla.org/index.html>



Alert details: <http://alerts.telemetry.mozilla.org/index.html#/detectors/1/metrics/677/alerts/?from=2015-09-18&to=2015-09-18>

Changeset for 20150918030202: <https://hg.mozilla.org/mozilla-central/pushloghtml?fromchange=e7d613b3bcfe1e865378bfac37de64560d1234ec&tochange=11dc79e232110ba6de5179e46dfbda77b52a88c3>

REFERENCES

- <http://robertovitillo.com/2015/07/02/telemetry-metrics-roll-ups/>
- <https://anthony-zhang.me/blog/telemetry-demystified/>
- <http://robertovitillo.com/2014/07/28/regression-detection-for-telemetry-histograms/>
- <https://telemetry.mozilla.org/>
- <https://metrics.services.mozilla.com/>

ADDING A PROBE


- Telemetry histograms are the preferred way to track numeric measurements such as timings
- There are various types of histograms [1]:
 - flag, e.g. FXA_CONFIGURED
 - boolean, e.g. E10S_WINDOW
 - count, e.g. CONTENT_DOCUMENTS_DESTROYED
 - enumerated, e.g. DEVICE_RESET_REASON
 - linear, e.g. GC_MAX_PAUSE_MS
 - exponential, e.g. GC_MARK_MS

[1] https://developer.mozilla.org/en-US/docs/Mozilla/Performance/Adding_a_new_Telemetry_probe


Keyed histograms are collections of one of the histogram types, indexed by a string key, e.g. SEARCH_COUNTS

SEARCH_COUNTS distribution for Firefox Desktop nightly 43, on any OS (67) any architecture (3) with any e10s setting (2) as any process type (2) and compare by none


google.searchbar




google.urlbar



google.abouthome



google.newtab



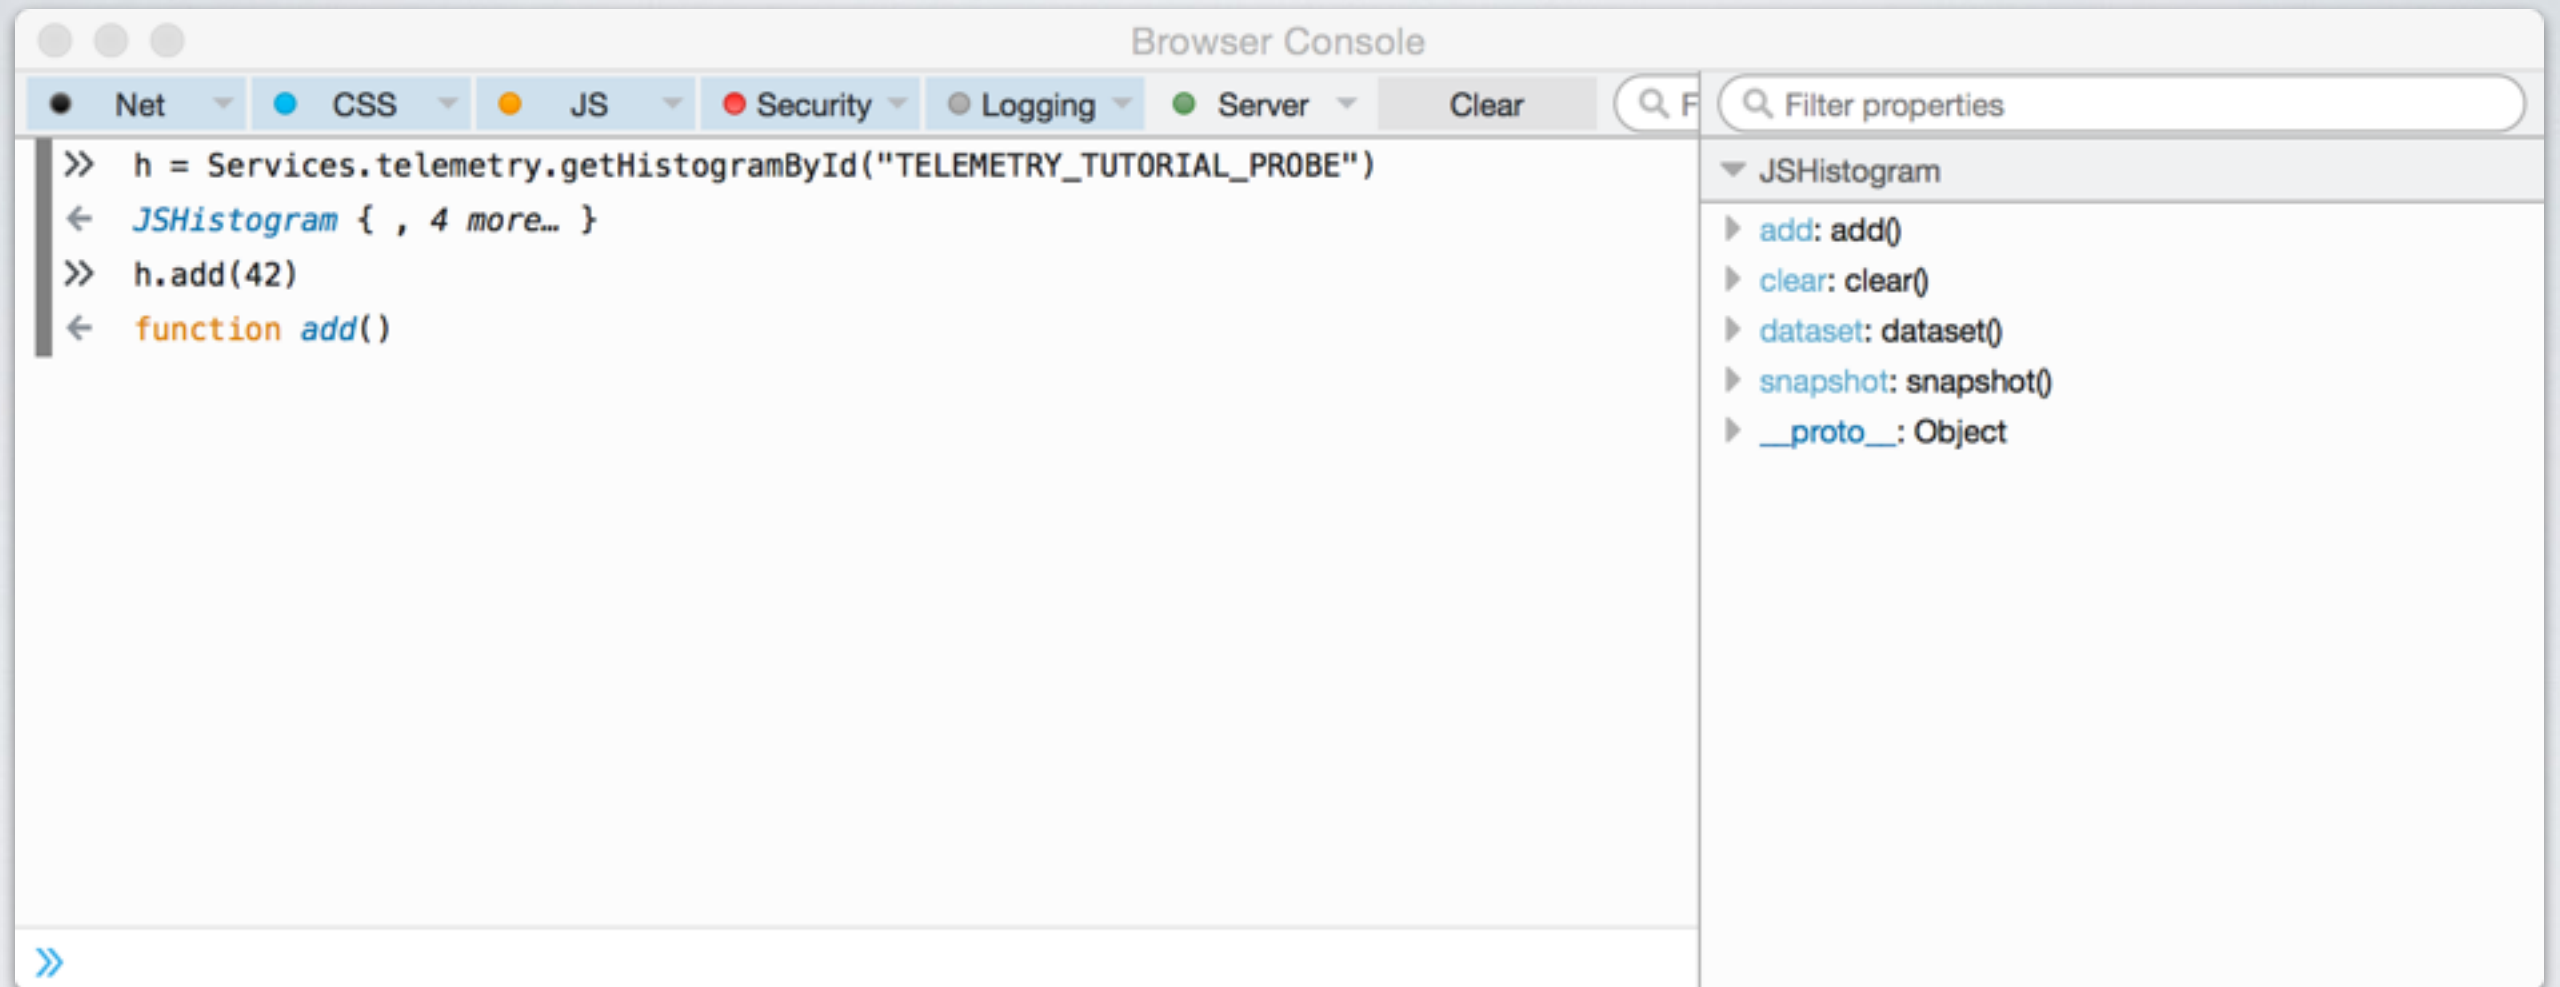
Declaring a Histogram

toolkit/components/telemetry/Histograms.json

```
"TELEMETRY_TUTORIAL_PROBE": {  
  "alert_emails": ["rvitillo@mozilla.com"],  
  "expires_in_version": "50",  
  "kind": "exponential",  
  "high": "10000",  
  "n_buckets": 50,  
  "description": "Telemetry tutorial probe (ms)"  
}
```

./mach build toolkit/components/telemetry

Accumulating data



REFERENCES

- https://developer.mozilla.org/en-US/docs/Mozilla/Performance/Adding_a_new_Telemetry_probe

WHAT IS A PING?

- A Telemetry ping [1] is the data that we send to our Telemetry servers.
- Stored on the client as a JSON object
- Contains common information to all pings and a payload specific to a certain ping types.

[1] https://ci.mozilla.org/job/mozilla-central-docs/Tree_Documentation/toolkit/components/telemetry/telemetry/pings.html

Ping Types

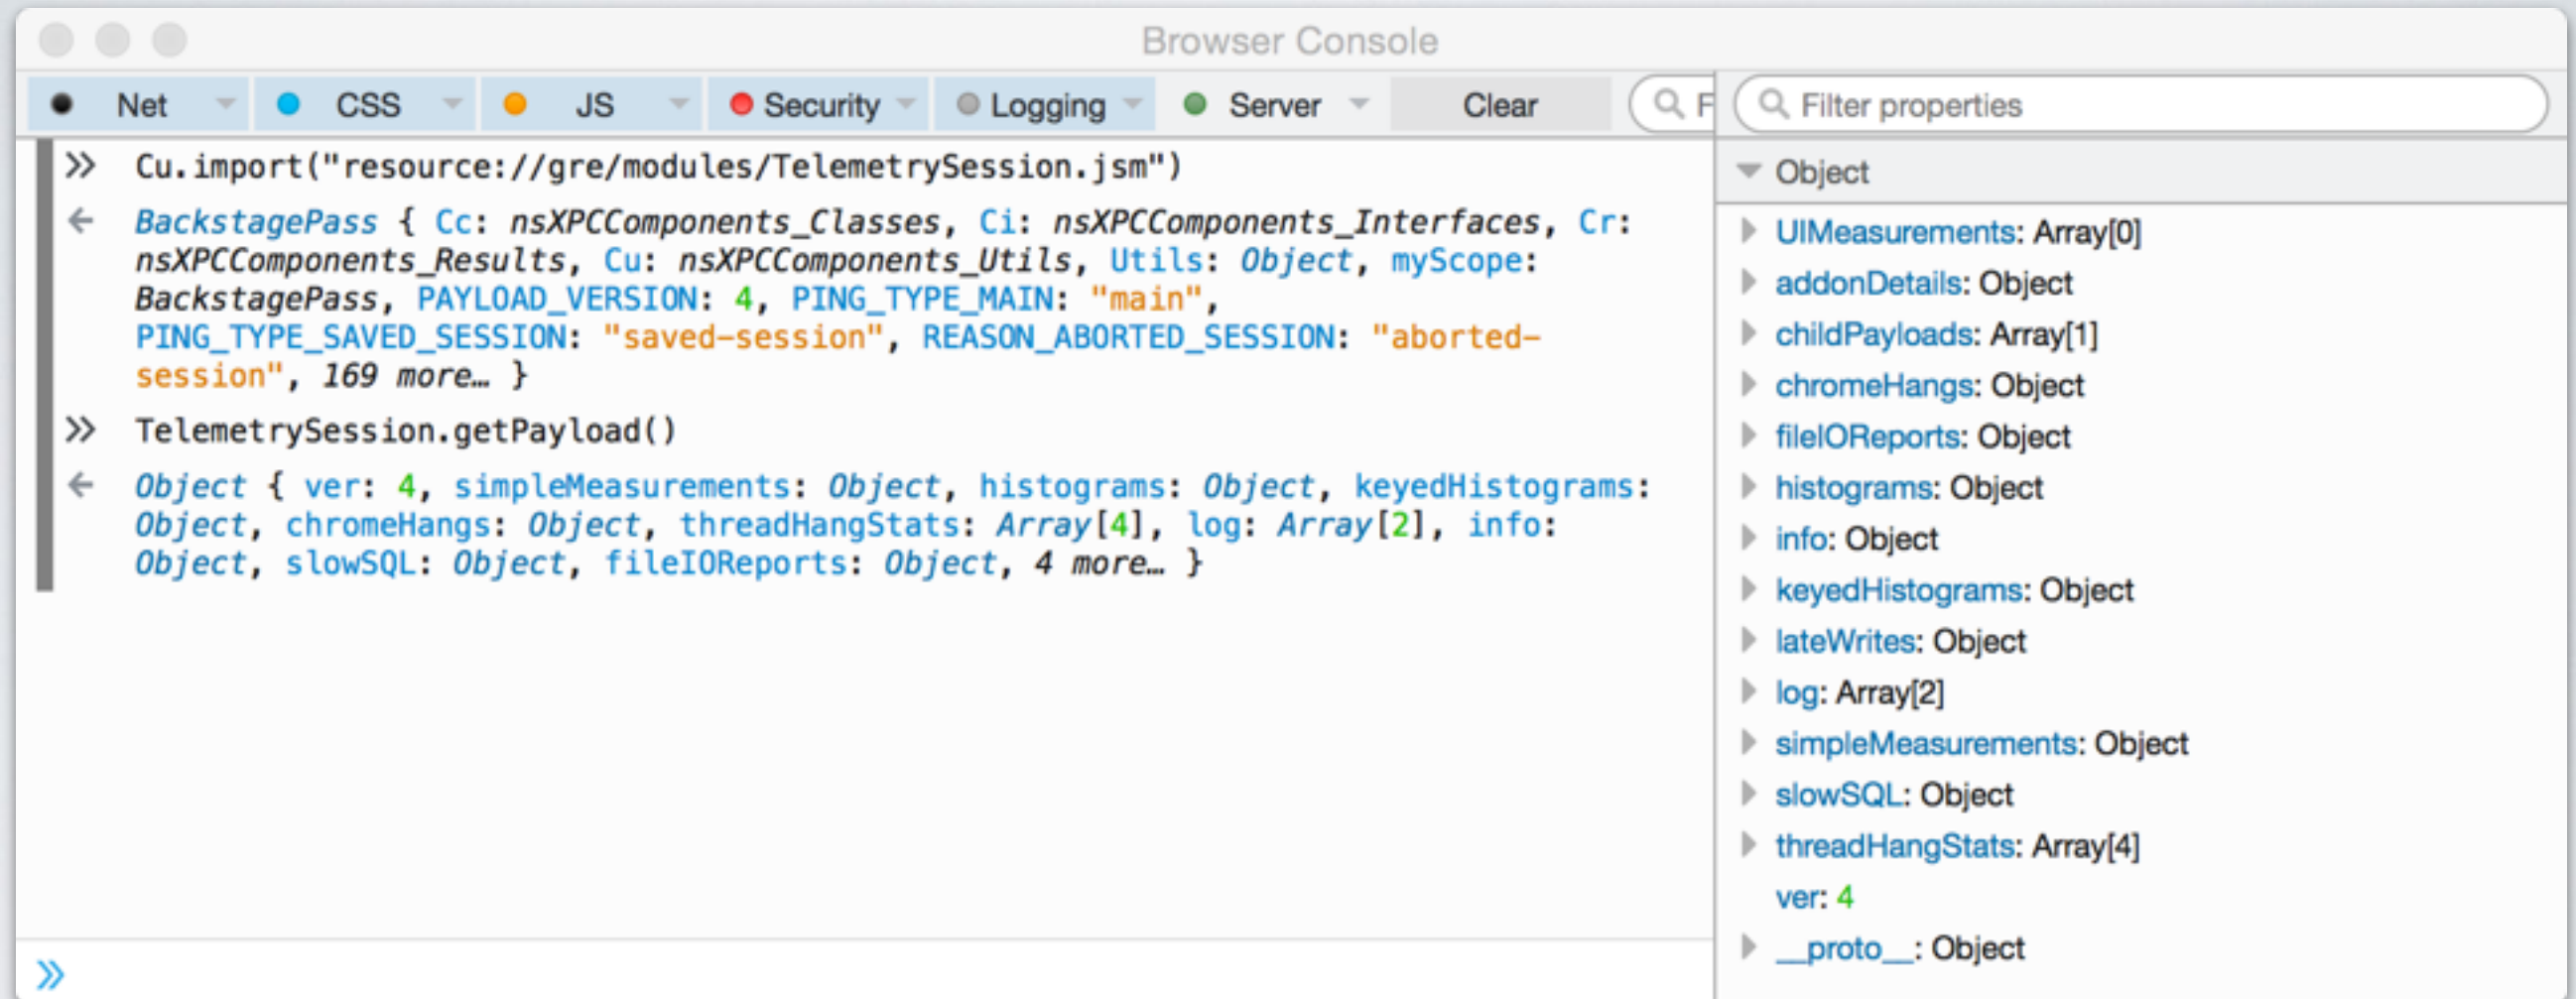
- Main ping, contains most of the measurements that are used to track the performance and health of Firefox in the wild
 - ▶ a ping is triggered by different scenarios, which is documented by the reason field (e.g. shutdown)
 - ▶ reasons lead to a session split, initiating a new sub-session; important measurements are reset for those subsessions
- Crash ping, captured after the main Firefox process crashes


```
{
  type: <string>, // "main", "activation", "deletion", "saved-session", ...
  id: <UUID>, // a UUID that identifies this ping
  creationDate: <ISO date>, // the date the ping was generated
  version: <number>, // the version of the ping format, currently 4

  application: {
    architecture: <string>, // build architecture, e.g. x86
    buildId: <string>, // "20141126041045"
    name: <string>, // "Firefox"
    version: <string>, // "35.0"
    vendor: <string>, // "Mozilla"
    platformVersion: <string>, // "35.0"
    xpcomAbi: <string>, // e.g. "x86-msvc"
    channel: <string>, // "beta"
  },

  clientId: <UUID>, // optional
  environment: { ... }, // optional, not all pings contain the environment
  payload: { ... }, // the actual payload data for this ping type
}
```

Main Ping Payload



```
>> Cu.import("resource://gre/modules/TelemetrySession.jsm")
< BackstagePass { Cc: nsXPCComponents_Classes, Ci: nsXPCComponents_Interfaces, Cr: nsXPCComponents_Results, Cu: nsXPCComponents_Utils, Utils: Object, myScope: BackstagePass, PAYLOAD_VERSION: 4, PING_TYPE_MAIN: "main", PING_TYPE_SAVED_SESSION: "saved-session", REASON_ABORTED_SESSION: "aborted-session", 169 more... }

>> TelemetrySession.getPayload()
< Object { ver: 4, simpleMeasurements: Object, histograms: Object, keyedHistograms: Object, chromeHangs: Object, threadHangStats: Array[4], log: Array[2], info: Object, slowSQL: Object, fileIOReports: Object, 4 more... }
```

Browser Console

Net CSS JS Security Logging Server Clear

Filter properties

Object

- UIMeasurements: Array[0]
- addonDetails: Object
- childPayloads: Array[1]
- chromeHangs: Object
- fileIOReports: Object
- histograms: Object
- info: Object
- keyedHistograms: Object
- lateWrites: Object
- log: Array[2]
- simpleMeasurements: Object
- slowSQL: Object
- threadHangStats: Array[4]
- ver: 4
- __proto__: Object

`Cu.import("resource://gre/modules/TelemetrySession.jsm")`

EIOs Caveat

The screenshot shows a browser console window titled "Browser Console". It has tabs for "Net", "CSS", "JS", "Security", "Logging", and "Server". The "JS" tab is selected. The console shows the following sequence of events:

- A script is imported: `Cu.import("resource://gre/modules/TelemetrySession.jsm")`.
- A function `BackstagePass` is called, returning an object with various properties like `nsXPCComponents_Classes`, `nsXPCComponents_Interfaces`, `nsXPCComponents_Results`, `nsXPCComponents_Utils`, `Utils`, `myScope`, `PAYLOAD_VERSION`, `PING_TYPE_MAIN`, `PING_TYPE_SAVED_SESSION`, and `REASON_ABORTED_SESSION`.
- The `TelemetrySession.requestChildPayloads()` method is called, returning `undefined`.
- The `TelemetrySession.getPayload()` method is called, returning an object with properties like `ver`, `simpleMeasurements`, `histograms`, `keyedHistograms`, `chromeHangs`, `threadHangStats`, `log`, `info`, `slowSQL`, and `fileIOReports`.

The right side of the console shows a detailed view of the object returned by `TelemetrySession.getPayload()`. It is an `Object` with the following properties:

- `UIMeasurements`: `Array[0]`
- `addonDetails`: `Object`
- `childPayloads`: `Array[1]`
 - `0`: `Object`
 - `chromeHangs`: `Object`
 - `histograms`: `Object`
 - `keyedHistograms`: `Object`
 - `log`: `Array[0]`
 - `simpleMeasurements`: `Object`
 - `threadHangStats`: `Array[3]`
 - `ver`: `4`
 - `__proto__`: `Object`
 - `length`: `1`
 - `__proto__`: `Array[0]`
- `chromeHangs`: `Object`
- `fileIOReports`: `Object`

Environment

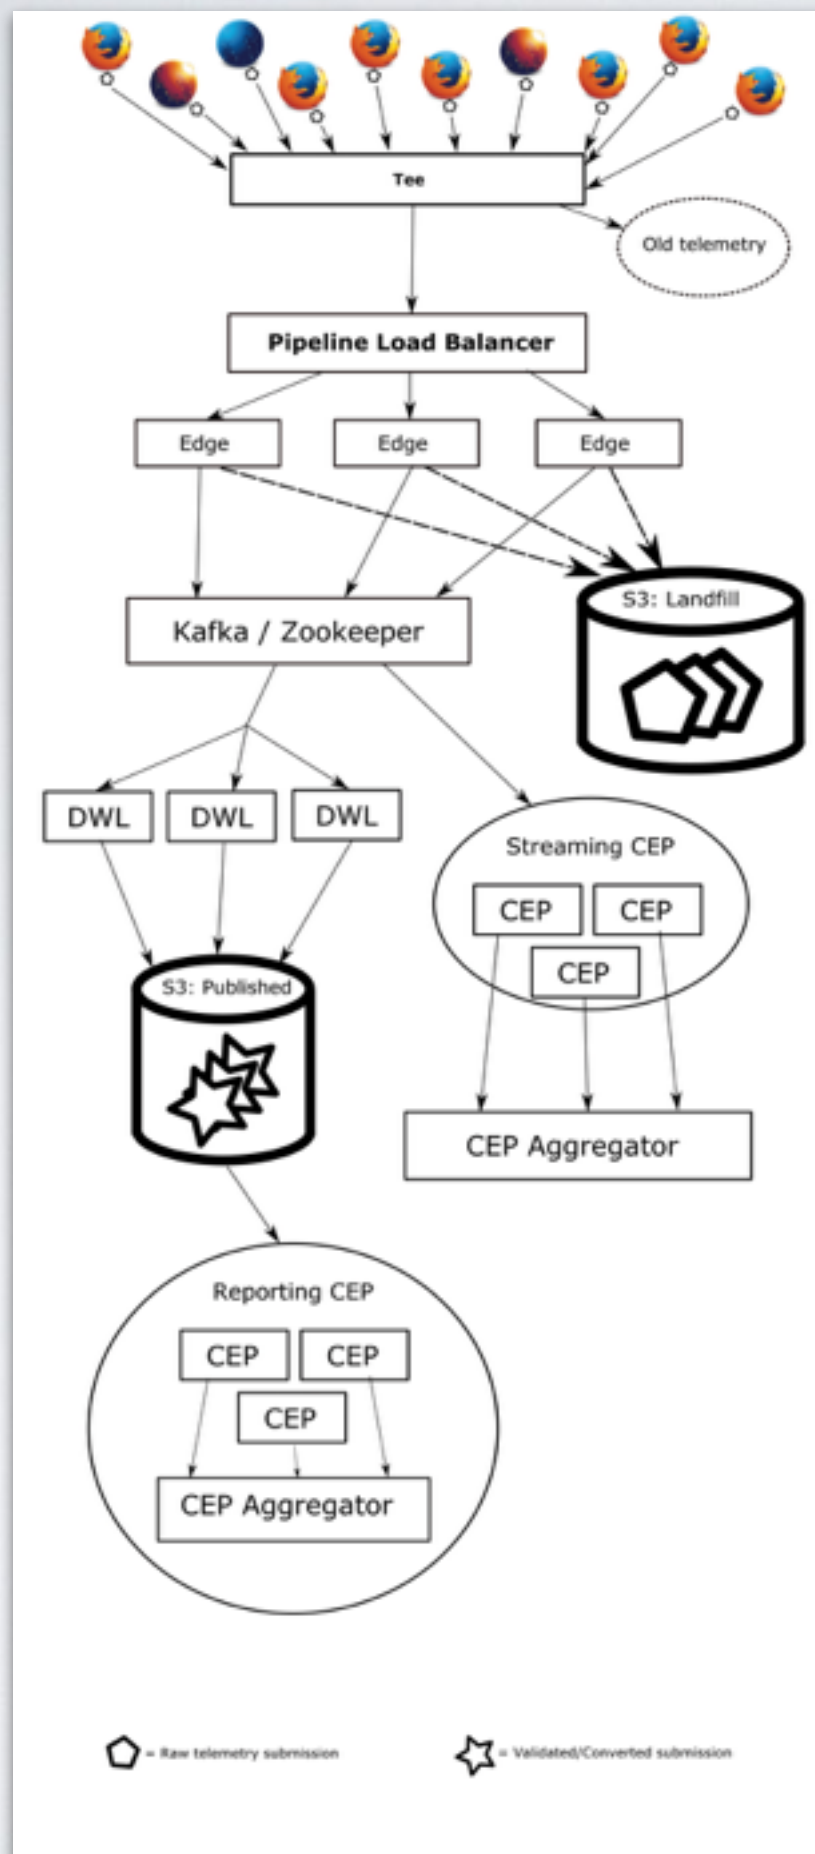
- Consists of data that is expected to be characteristic for performance and other behaviour and not expected to change too often (e.g. number of CPU cores)
- Changes to most of these data points are detected and lead to a session split in the “main” ping

REFERENCES

- https://ci.mozilla.org/job/mozilla-central-docs/Tree_Documentation/toolkit/components/telemetry/telemetry/pings.html

DATA PIPELINE

- The data pipeline ingests, transforms, stores and analyzes data
- The endpoint is an HTTP server that:
 - ▶ listens for POST/PUT from Firefox
 - ▶ does some decoding / preprocessing
 - ▶ makes the data available for streaming analyses (Heka)
 - ▶ archives the data to S3 for further offline analyses (Spark)



CEP = Complex Event Processor (basically streaming analysis or reporting)

DWL = Data Warehouse Loader

Landfill = Shorthand for "write-mostly store for backup / recovery purposes"

REFERENCES

- <https://wiki.mozilla.org/CloudServices/DataPipeline>
- <https://github.com/mozilla-services/data-pipeline>
- <https://github.com/mozilla-services/heka>
- <https://github.com/apache/spark/>

Break Time!

Regroup in a few minutes

OFFLINE PROCESSING

IPYTHON, PANDAS & SPARK

<https://github.com/vitillo/telemetry-onboarding/tree/master/notebooks>

EXPERIMENTS

- Telemetry Experiments is a feature that allows Firefox to automatically download and run specially-designed restartless addons based on certain conditions
- Currently only available in desktop versions of Firefox
- Experiments can be targeted at various populations by specifying conditions in an experiment manifest
- Before building an experiment, contact a data-collection peer; additional privacy or security review might be required
- Product approval is required to run an experiment

1. Setup staging server [1], which Firefox is going to contact to install the current experiment
2. Write restart-less add-on [2]
 - `addon.js`, when `startup()` is called, it must manually inject its user interface and other behavior into the application; when `shutdown()` is called, it must remove anything that it has added to the application
 - `install.rdf`, used to determine information about an add-on as it is being installed (e.g. author, version, etc.)
3. `manifest.json`, the experiment manifest which targets a specific population

[1] https://wiki.mozilla.org/QA/Telemetry/Installing_Staging_Server

[2] https://developer.mozilla.org/en-US/Add-ons/Bootstrapped_extensions

Flags required for testing

- `experiments.force-sample-value = "0.0"`
- `experiments.logging.level = 0`
- `experiments.manifest.cert.checkAttributes = false`
- `experiments.manifest.uri = "http://localhost:8000/firefox-manifest.json"`
- `xpinstall.signatures.required = false`


```

1 let {classes: Cc, interfaces: Ci, utils: Cu} = Components;
2
3 Cu.import("resource:///modules/experiments/Experiments.jsm");
4 Cu.import("resource:///gre/modules/Task.jsm");
5 Cu.import("resource:///gre/modules/Preferences.jsm");
6
7 var gStarted = false;
8
9 const kSELF_ID = "flash-protectedmode-beta35@experiments.mozilla.org";
10
11 function startup() {
12     // Seems startup() function is launched twice after install, we're
13     // unsure why so far. We only want it to run once.
14     if (gStarted) {
15         return;
16     }
17     gStarted = true;
18
19     Task.spawn(function*() {
20         let branch = yield Experiments.instance().getExperimentBranch(kSELF_ID);
21         switch (branch) {
22             case "control":
23                 return;
24             case null:
25                 let r = (Math.random() >= 0.5);
26                 if (!r) {
27                     yield Experiments.instance().setExperimentBranch(kSELF_ID, "control");
28                     return;
29                 }
30                 yield Experiments.instance().setExperimentBranch(kSELF_ID, "experiment");
31                 // FALL THROUGH
32             case "experiment":
33                 let defaultPrefs = new Preferences({defaultBranch: true});
34                 defaultPrefs.set("dom.ipc.plugins.flash.disable-protected-mode", true);
35                 return;
36             default:
37                 throw new Error("Unexpected experiment branch: " + branch);
38         }
39     }).then(
40         function() {
41         },
42         function(e) {
43             Cu.reportError("Got error during bootstrap startup: " + e);
44         });
45 }
46
47 function shutdown() {
48     let defaultPrefs = new Preferences({defaultBranch: true});
49     defaultPrefs.set("dom.ipc.plugins.flash.disable-protected-mode", false);
50 }

```



```
1 {
2   "publish"      : true,
3   "priority"     : 2,
4   "name"         : "Flash Protected-Mode Testing",
5   "description"  : "Measuring the effect of Flash protected mode on crashes, hangs, and other browser jank.",
6   "info"         : "<p><a href=\"https://bugzilla.mozilla.org/show_bug.cgi?id=1110215\">Related bug</a></p>",
7   "manifest"     : {
8     "id"          : "flash-protectedmode-beta35@experiments.mozilla.org",
9     "startTime"   : 1418601600,
10    "endTime"     : 1421280000,
11    "maxActiveSeconds" : 2764800,
12    "appName"     : ["Firefox"],
13    "channel"     : ["beta"],
14    "os"          : ["WINNT"],
15    "minVersion"  : "35.0",
16    "minBuildID"  : "20141215000000",
17    "maxVersion"  : "37.*",
18    "sample"      : 0.1,
19    "disabled"    : true
20  }
21 }
```

Bug 1111791 - Telemetry report: effect of the Flash protected-mode experiment

https://bugzilla.mozilla.org/show_bug.cgi?id=1111791

REFERENCES


- <https://wiki.mozilla.org/Telemetry/Experiments>
- https://developer.mozilla.org/en-US/Add-ons/Bootstrapped_extensions
- https://wiki.mozilla.org/QA/Telemetry#Telemetry_Experiments.2FFHR_Documentation
- <http://codefirefox.com/video/install-telmetry-experiment>
- <http://hg.mozilla.org/webtools/telemetry-experiment-server/file/tip/experiments>
- https://bugzilla.mozilla.org/show_bug.cgi?id=1110215
- https://bugzilla.mozilla.org/show_bug.cgi?id=1111791

STATS

USE REPRESENTATIVE SAMPLES




USE SUFFICIENT DATA




A 95% confidence interval indicates that 19 out of 20 samples (95%) from the same population will produce confidence intervals that contain the population parameter.

BEWARE OF PSEUDOREPLICATION



CORRELATION IS NOT CAUSATION




USE THE RIGHT AVERAGE

- Mean: add up the values, and divide by the number of values
- Median: the median is the 50th percentile; half the values are higher than the median, and half are lower
- Geometric mean: compute the logarithm of all values, compute the mean of the logarithms, and then take the antilog
 - It is a better measure of central tendency when data follow a lognormal distribution (long tail).

REGRESSION TOWARD THE MEAN

Regression / Reversion to the Mean



If the red line is the average (mean), the trend oscillates up and down around it.

The trend moves above the average, then regresses back toward it, often overshooting up and down.

MAKE YOUR ANALYSES REPRODUCIBLE

```
In [1]: import ujson as json
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import plotly.plotly as py

from moztelemetry import get_pings, get_pings_properties, get_one_ping_per_client, get_clients_history

%pylab inline
```

Populating the interactive namespace from numpy and matplotlib

```
In [7]: pings = get_pings(sc, app="Firefox", channel="release", submission_date="20150928", build_id="20150917150946")
```

```
In [8]: def telemetry_enabled(ping):
        return ping.get("environment", {}).get("settings", {}).get("telemetryEnabled", False)
```

```
In [*]: pings.count()
```

```
Out[9]: 1242836
```


REFERENCES

- <http://www.statisticsonewrong.com/data-analysis.html>
- <http://www.amazon.com/How-Lie-Statistics-Darrell-Huff/dp/0393310728>
- <http://www.slideshare.net/RobertoAgostinoVitali/all-you-need-to-know-about-statistics>

PRIVACY POLICY

1. No surprises: use and share information in a way that is transparent and benefits the user
2. User control: develop products and advocate for best practices that put users in control of their data and online experiences
3. Limited data: collect what we need, de-identify where we can and delete when no longer necessary
4. Sensible settings: design for a thoughtful balance of safety and user experience
5. Defense in depths: maintain multi-layered security controls and practices, many of which are publicly verifiable

DATA COLLECTION POLICY

- When proposing a new measurement or data system, consider the requirements and the necessary data properties, e.g:
 - ▶ is it necessary to take a measurement from all users? Or is it sufficient to measure only prerelease users?
 - ▶ is it desirable to track data changes over time? With what frequency? With what latency?
- For every new measurement, even a simple new Telemetry probe, please request approval by setting the feedback flag for the module owner or a peer.
 - ▶ Owner: Benjamin Smedberg
 - ▶ Peers: Vladan Djerić, Ally Naaghtgeboren

